

## Bonus Questions:

1. How would you optimize the ``get_jobs`` function?
  - a. This could be done using better data\_structure , instead of iterating the entire list.
  - b. When using a data structure that searches pending jobs in  $O(1)$ .  
If we are using an external database, we can index jobs based on status.
  - c. To obtain a fair scheduling we can also use Multi-level-feedback Queue.
2. Are there tools that you would use instead of writing this script to manage the job scheduling? How would the entire solution change to adopt them?
  - a. Cron Job - Linux has a cron job utility where we can schedule different tasks, also Kubernetes cron is an excellent choice for such a system
3. What would you do differently if the job was CPU-bound rather than IO-bound? Particularly since Python is not a parallel language (i.e. GIL).
  - a. Since Python is not a parallel language multi-threading does not make sense in CPU Bound processes. We can instead use Multi-Processors where each core has its own GIL.  
We would also need to take care of how graceful shutdowns are handled on each core.
4. How should someone deploying a scheduler-powered job determine their value for ``--max-concurrent-jobs``?
  - a. It should depend on number of cores available to run CPU-Bound process in a multi-processing environment.